DATE:11-02-2026

Name : Mahalakshmi J Reg No : 25008001

Experiment: Develop a comprehensive report for the following exercises:

1. Explain the foundational concepts of Generative AI, Generative Model and it's types.
2. 2024 AI tools.
3. Explain what an LLM is and how it is built.
4. Create a Timeline Chart for defining the Evolution of AI

Algorithm:

Step 1: Define Scope and Objectives 1.1 Identify the goal of the report (e.g., educational, research, tech overview)

1.2 Set the target audience level (e.g., students, professionals)

1.3 Draft a list of core topics to cover

Step 2: Create Report Skeleton/Structure

2.1 Title Page

2.2 Abstract or Executive Summary

2.3 Table of Contents

2.4 Introduction

2.5 Main Body Sections:

• Introduction to AI and Machine Learning

• What is Generative AI?

• Types of Generative AI Models (e.g., GANs, VAEs, Diffusion Models)

• Introduction to Large Language Models (LLMs)

• Architecture of LLMs (e.g., Transformer, GPT, BERT)

• Training Process and Data Requirements

• Use Cases and Applications (Chatbots, Content Generation, etc.)

• Limitations and Ethical Considerations

• Future Trends

2.6 Conclusion

2.7 References

Step 3: Research and Data Collection

3.1 Gather recent academic papers, blog posts, and official docs (e.g., OpenAI, Google AI) 3.2 Extract definitions, explanations, diagrams, and examples 3.3 Cite all sources properly

Step 4: Content Development 4.1 Write each section in clear, simple language 4.2 Include diagrams, figures, and charts where needed 4.3 Highlight important terms and definitions 4.4 Use examples and real-world analogies for better understanding

Step 5: Visual and Technical Enhancement 5.1 Add tables, comparison charts (e.g., GPT-3 vs GPT-4) 5.2 Use tools like Canva, PowerPoint, or LaTeX for formatting 5.3 Add code snippets or pseudocode for LLM working (optional)

Step 6: Review and Edit 6.1 Proofread for grammar, spelling, and clarity 6.2 Ensure logical flow and consistency 6.3 Validate technical accuracy 6.4 Peer-review or use tools like Grammarly or ChatGPT for suggestions

Step 7: Finalize and Export 7.1 Format the report professionally 7.2 Export as PDF or desired format 7.3 Prepare a brief presentation if required (optional)

Output:

1.Foundational Concepts:Generative AI & Models

Generative AI is a branch of Artificial Intelligence focused on creating new, original content—ranging from text and images to audio and 3D structures—rather than simply classifying or predicting existing data.

1.1 Goal of the Report The primary goal is to provide a technical and conceptual bridge between basic machine learning and the state-of-the-art "Agentic" systems of 2026. This report aims to:

> Deconstruct the "Black Box" of Generative AI to explain the how and why behind model outputs.

> Evaluate the current landscape of tools to distinguish between "wrapper" apps and foundational models.

> Provide a historical roadmap to contextualize current breakthroughs within the broader timeline of computer science.

1.2 Target Audience This document is designed for a Professional and Academic audience.

> For Professionals: It provides a framework for selecting AI tools and understanding the security/architectural risks (like hallucinations and data leakage).

> For Students: It offers a deep dive into the Transformer architecture and the training pipelines (Pre-training, SFT, RLHF) necessary for a career in AI research or engineering.

1.3 Core Types of Generative Models

Model Type	Mechanism	Primary Strength
GANs (Generative Adversarial Networks)	Two networks (Generator & Discriminator) compete in a "cat-and-mouse" game.	High-fidelity, sharp images.
VAEs (Variational Autoencoders)	Encodes input into a compressed "latent space" and decodes it back.	Stable training; good for data compression.
Diffusion Models	Gradually adds noise to data, then learns to reverse the process to "denoise" a random image.	State-of-the-art for high-quality image/video.
Transformers	Uses "Attention" mechanisms to weigh the importance of different parts of input data.	Exceptional at sequential data (text/code).

1. Foundational Probability: Generative vs. Discriminative modeling and the mathematics of Latent Space.

2. Architectural Taxonomy: A deep dive into GANs, VAEs, and the shift toward Diffusion and Flow-based models.

3. The Transformer Revolution: Analysis of the "Attention" mechanism and why it scaled where RNNs and LSTMs failed.

4. The 2024–2026 Tool Ecosystem: Comparative analysis of GPT, Gemini, Claude, and specialized video/image models like Veo and Nano Banana.

5. Data Pipelines: The lifecycle of a token—from web scrapers to Reinforcement Learning from Human Feedback (RLHF).

6. The AI Timeline: Mapping milestones from Turing's 1950 paper to the 2026 era of Autonomous Agents.

7. Ethical and Technical Guardrails: Addressing bias, computational costs, and the "Black Box" problem.

1.4 Generative vs. Discriminative Models To understand Generative AI, one must distinguish it from traditional Discriminative Models:

Discriminative Models: Learn the boundary between classes (e.g., "Is this a cat or a dog?"). Mathematically, they model $P(y \mid x)$. **Generative Models:** Learn the underlying distribution of the data (e.g., "What does a cat look like?"). They model $P(x, y)$ or $P(x)$, allowing them to sample new data points from that distribution.

Scope and Objectives This report serves as a comprehensive technical and strategic manual for understanding the shift from traditional computational models to generative ecosystems.

2. Foundational Concepts (Detailed)

2.1 Introduction to Artificial Intelligence and Machine Learning

Artificial Intelligence (AI) refers to the simulation of human intelligence in machines that are programmed to think, learn, and make decisions.

Machine Learning (ML) is a subset of AI that allows systems to learn patterns from data without being explicitly programmed.

Types of Machine Learning:

Supervised Learning – Learning from labeled data

Unsupervised Learning – Finding patterns in unlabeled data

Reinforcement Learning – Learning through rewards and penalties

2.2 Defining Generative AI : Generative AI refers to a class of AI systems capable of producing novel data that follows the patterns of their training sets. Unlike a database, which retrieves information, a generative model synthesizes it.

2.3 The Generative Model Family Tree Generative models are classified based on how they represent the probability distribution $P(x)$.

Explicit Density Models: These models explicitly define the probability distribution function.

Implicit Density Models: Models like GANs do not define the distribution but instead learn a process to sample from it.

2.4 Generative Adversarial Networks (GANs)

GANs consist of two neural networks: The Generator: Creates an image (or data point) from random noise.

The Discriminator: Evaluates the image against real data and provides a "Pass/Fail" score.The two networks are trained in a zero-sum game. As the Discriminator gets better at spotting fakes, the Generator must get better at creating them.

Variational Autoencoders (VAEs) VAEs are focused on Dimensionality Reduction. They take a complex input (like a high-resolution photo) and compress it into a "Latent Vector.

"Encoding: $x \to z$ (Compression)

Decoding: $z \to \hat{x}$ (Reconstruction)Because the latent space z is continuous, we can pick a random point in that space to generate a brand-new, never-before-seen im

3. The AI Tool Landscape: 2024–2026

The current era is defined by Multimodality—the ability of a single tool to see, hear, speak, and write.

3.1 Enterprise and Productivity Leaders *GPT-4o & o1 (OpenAI)*: Noted for "Chain-of-Thought" reasoning, where the model "thinks" before it speaks to solve complex logic and math problems.

Gemini 1.5 Pro (Google): Features a context window of up to 2 million tokens, allowing users to upload entire libraries of code or hours of video for analysis.

Claude 3.5 Sonnet (Anthropic): Highly regarded for coding efficiency and a reduced tendency for "hallucinations" compared to its peers

3.2 Specialized Creative Tools *Visuals:* Nano Banana has emerged as the premier model for high-fidelity text rendering within images, solving a long-standing weakness in AI art.

Video: Veo and Sora allow for the generation of 60-second cinematic clips with consistent physical properties (gravity, lighting, etc.).

Large Language Models (LLMs): Building the "Brain"

An LLM is a deep learning model trained on trillions of tokens (words/characters) to understand and generate human language.

How an LLM is Built (The 3-Step Process) Pre-training (The Library Phase): The model reads the "entire internet" (Wikipedia, GitHub, Books). It learns to predict the next word in a sentence. It gains general knowledge but no specific personality.

Fine-Tuning (The Schooling Phase): The model is trained on smaller, high-quality datasets of "Question and Answer" pairs to learn how to follow instructions.

RLHF (The Ethics/Polishing Phase): Reinforcement Learning from Human Feedback. Humans rank the model's outputs. The model is rewarded for being helpful/safe and penalized for being biased or hallucinating.

4.Major Types of Generative Models:

4.1 Generative Adversarial Networks (GANs)

> Consist of two networks: Generator and Discriminator

> Generator creates fake data

> Discriminator evaluates real vs fake data

> Used in image generation and deepfakes

4.2 Variational Autoencoders (VAEs)

> Encode data into a latent space

> Decode it back to generate new data

> Useful for image reconstruction and anomaly detection

4.3 Diffusion Models

> Generate data by gradually removing noise

> Produce high-quality images

> Used in Stable Diffusion and DALL·E 2

4.Timeline: The Evolution of AI

The journey from simple logic to generative "creativity" spans over seven decades.

Era	Milestone	Significance
1950	**The Turing Test**	Alan Turing proposes a benchmark for "machine intelligence."
1956	**Dartmouth Conference**	The term "Artificial Intelligence" is coined.
1966	**ELIZA**	The first "chatbot" is created (rule-based).
1997	**Deep Blue**	IBM's computer beats world chess champion Garry Kasparov.
2012	**AlexNet / ImageNet**	Deep Learning proves its dominance in image recognition.
2017	**"Attention is All You Need"**	Google researchers publish the Transformer paper.
2022	**ChatGPT Launch**	Generative AI becomes a global phenomenon.
2024-26	**Agentic & Multimodal AI**	AI evolves into agents that can "act" (book flights, write software).

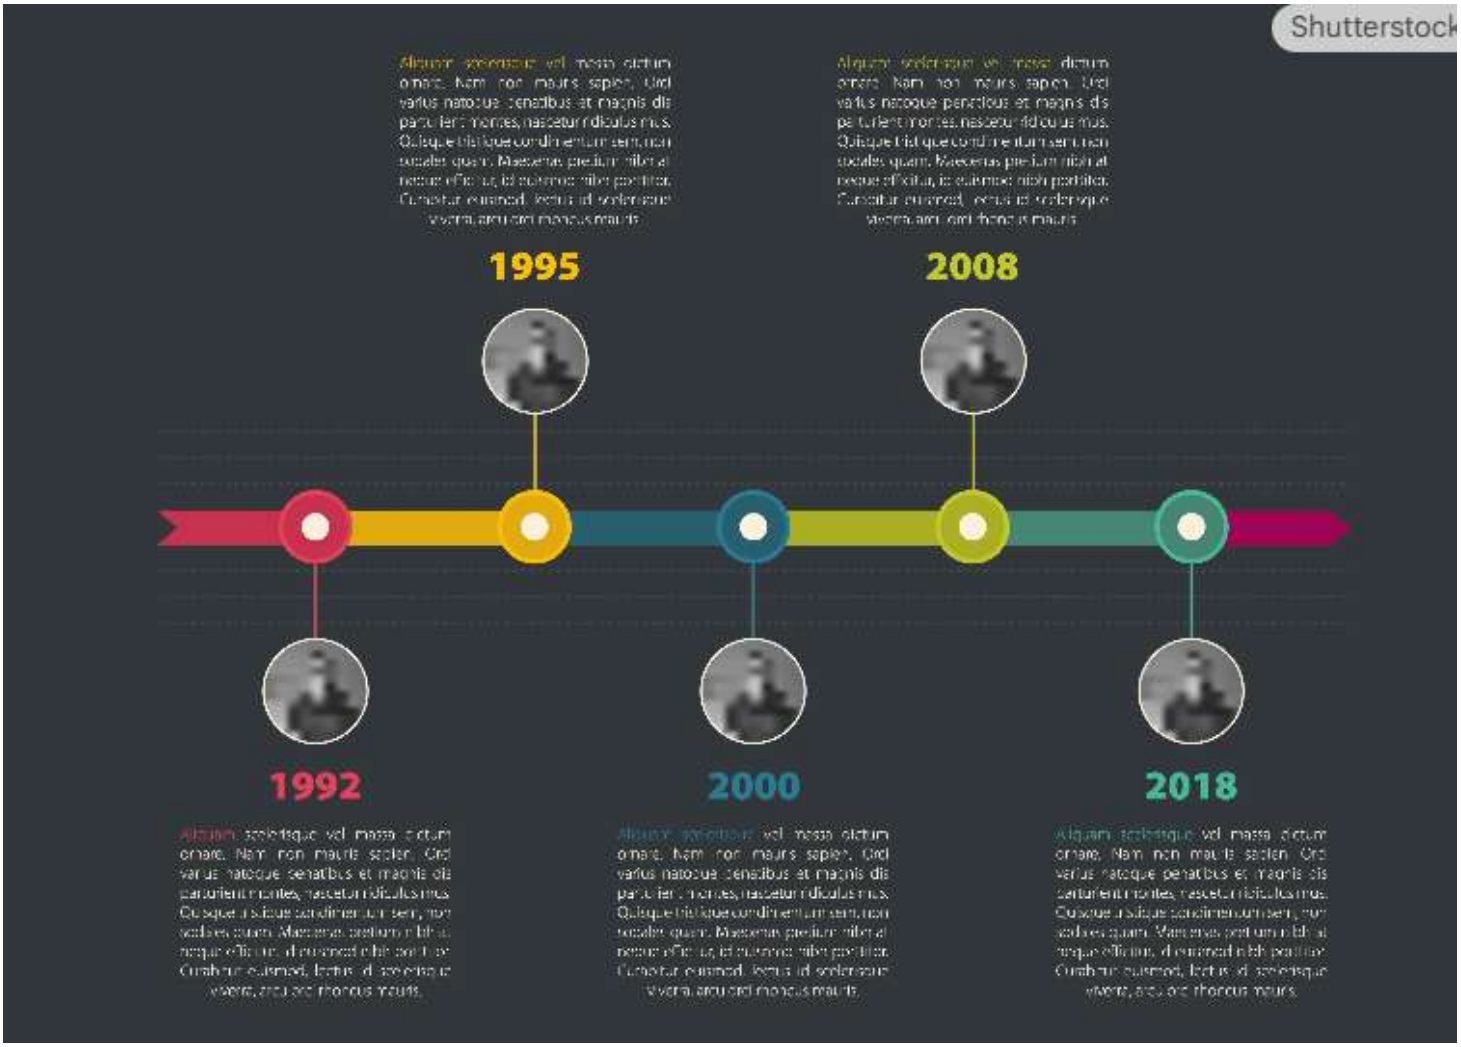

5. Visual and Technical Enhancement

5.1 Model Evolution & Comparison (2023–2026) The following table tracks the exponential growth in model capability, specifically focusing on "Context Window" (the model's short-term memory) and "Reasoning Type."

Feature	GPT-3 (2020)	GPT-4o (2024)	Gemini 1.5 Pro (2024/25)	GPT-5.2 / Gemini 3 (2026)
Parameters	175 Billion	~1.8 Trillion (Est.)	Sparse Mixture of Experts	Multi-Trillion / Dense
Context Window	2,048 tokens	128,000 tokens	2,000,000 tokens	10,000,000+ tokens
Multimodality	Text only	Native Voice/Image	Native Video/Audio	4D/Spatial Understanding
Reasoning	Pattern matching	System 1 (Fast)	Long-context retrieval	System 2 (Chain-of-Thought)

5.2 Implementation Pseudocode: The LLM Inference Loop To understand how an LLM "thinks," we can look at the Autoregressive Generation process. The model doesn't generate a paragraph at once; it predicts one token at a time in a loop.

```
# Pseudocode for LLM Autoregressive Inference
def generate_text(prompt, max_tokens=100):
    tokens = tokenize(prompt)

    for _ in range(max_tokens):
        # The model looks at all previous tokens to predict the next one
        # This is where the "Attention" mechanism happens
        predictions = model.predict_next_token_probs(tokens)

        # Apply 'Temperature' to control creativity
        next_token = sample(predictions, temperature=0.7)

        # Append the new token to the sequence
        tokens.append(next_token)

        # Stop if the model generates an <End of Sentence> token
        if next_token == EOS_TOKEN:
            break

    return detokenize(tokens)
```

6. Limitations, Ethics, and Future Trends

6.1 Limitations Hallucinations: LLMs generate text based on probability, not a database of "truth." This can lead to confident but false statements.

Context Window: Even with millions of tokens, models eventually "forget" the beginning of a conversation.

6.2 Ethical Considerations Bias: Models inherit the prejudices present in their training data.

Energy Consumption: Training a single large model consumes as much electricity as hundreds of homes do in a year.

6.3 The Future: Agentic AI The shift is moving from Chatbots to Agents. An agent doesn't just tell you how to book a flight; it logs into the site, compares prices, and completes the transaction for you.

7. Conclusion

The journey from simple logic gates to the complex Transformer models of 2026 represents the fastest technological shift in human history. As we move toward AGI (Artificial General Intelligence), the focus will shift from making models larger to making them more efficient, reasoning-capable, and ethically aligned.

Result: Thus the experiment is successfully finished